Full corpus ledger integrity audit

2026-10-21; crypto-autoresearcher main e20013ae20d2

Method. A full checkout held 138,890 tracked files. The repository's `tools/validate_ledger.py` ran with its existing prune-only baseline. It failed on **162 unsuppressed error lines in 49 records**. It also reported 1,210 grandfathered legacy errors and 543 advisories. None of these counts is a count of false mathematical or performance conclusions.

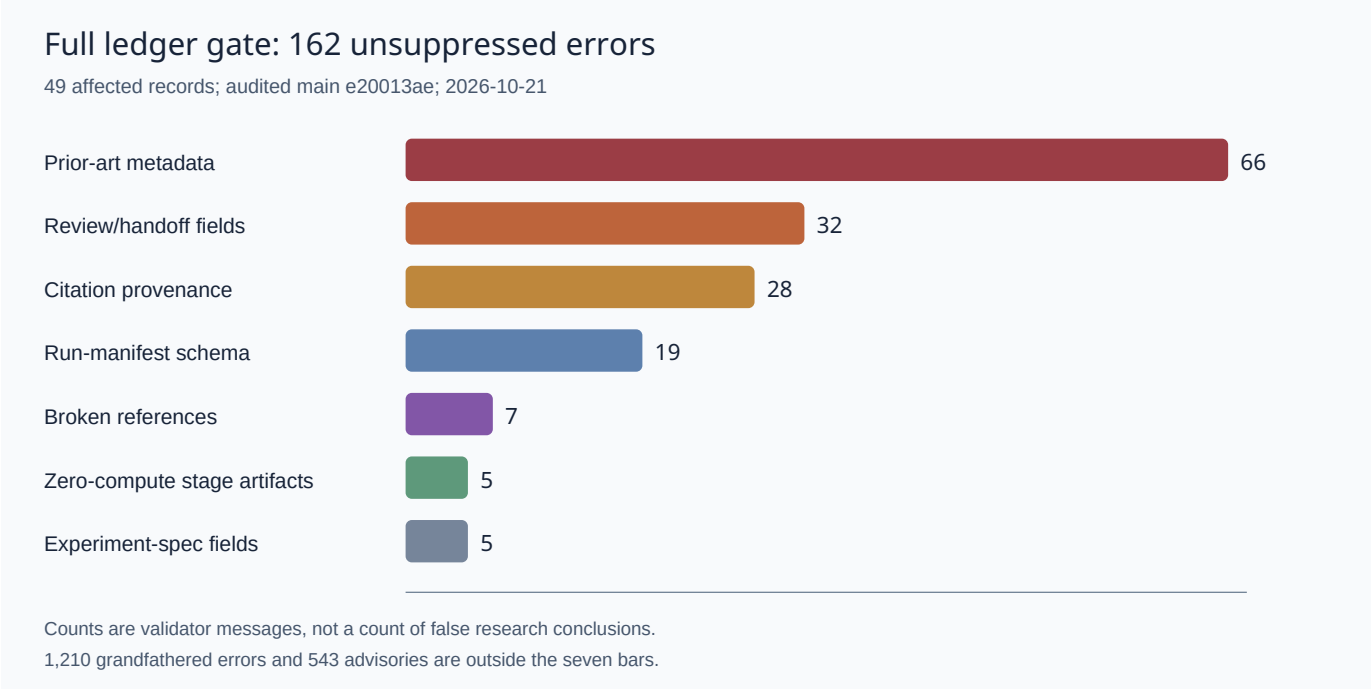

Consequence. The current ledger gate is red. The highest-volume classes concern prior-art metadata, citation provenance and blind review plans. Run-manifest and experiment-spec defects and broken references need record-by-record repair. Five absent companion artifacts belong to a declared zero-compute hand-derivation stage: its manifest says no program ran, so creating synthetic stdout or raw-result files would fabricate an execution. The classification or validator contract needs an honest repair.

Dashboard check. The complete-tree reader suite passed 145 tests. The local static build emitted 55,218 files (520.6 MB); its comparison payload carried all nine primary archived rows without a benchmark error. A direct check of 54 fields against the current unchanged primary archive found zero mismatches. The reader reports 24 unparseable historical source files: nine have hash-verified, parseable supersessions and fifteen have no registered replacement. The public Pages workflow is manually disabled; its live snapshot still names an October 11 commit. This build is local, not a deployment or a performance replay. The separate progress dashboard has no claim JSON inputs and therefore no computed frontiers.

Reader follow-on. The audit branch now exposes each of the nine hash-matched, parseable schema copies in a separate detail tab while retaining the original path, raw source, parse error and integrity warning. The updated reader suite passed 146 tests; a second complete local build kept 24 parse failures and all nine primary benchmark rows. Fifteen malformed originals have no registered copy. Hash matching is not semantic review, and the public site has not been refreshed. Browser interaction was not exercised locally because no browser runtime was available.

Boundary. The four newly supplied Lean targets and the complete 66-target machine gate are separate evidence. They do not clear these ledger errors, independent semantic review, or historical performance validation. No numerical graph changed because this audit produced no new timing datum. The exact validator log, classification script, source commit and follow-up conditions are in the accompanying Markdown report.