250-billion campaign launch

14 October 2026. Execution observations only: the campaign is running and the study is incomplete. No interim significance test or security verdict.

Measured: 150,994,944 checked point observations in 16 of 27,133 chunks; 589,824 window observations and 262,144 start draws. Zero failures / infrastructure stops. All completed chunks passed score replay and separate coordinate-frame checks. Snapshot: 2026-10-14T14:41:33.286299+00:00.

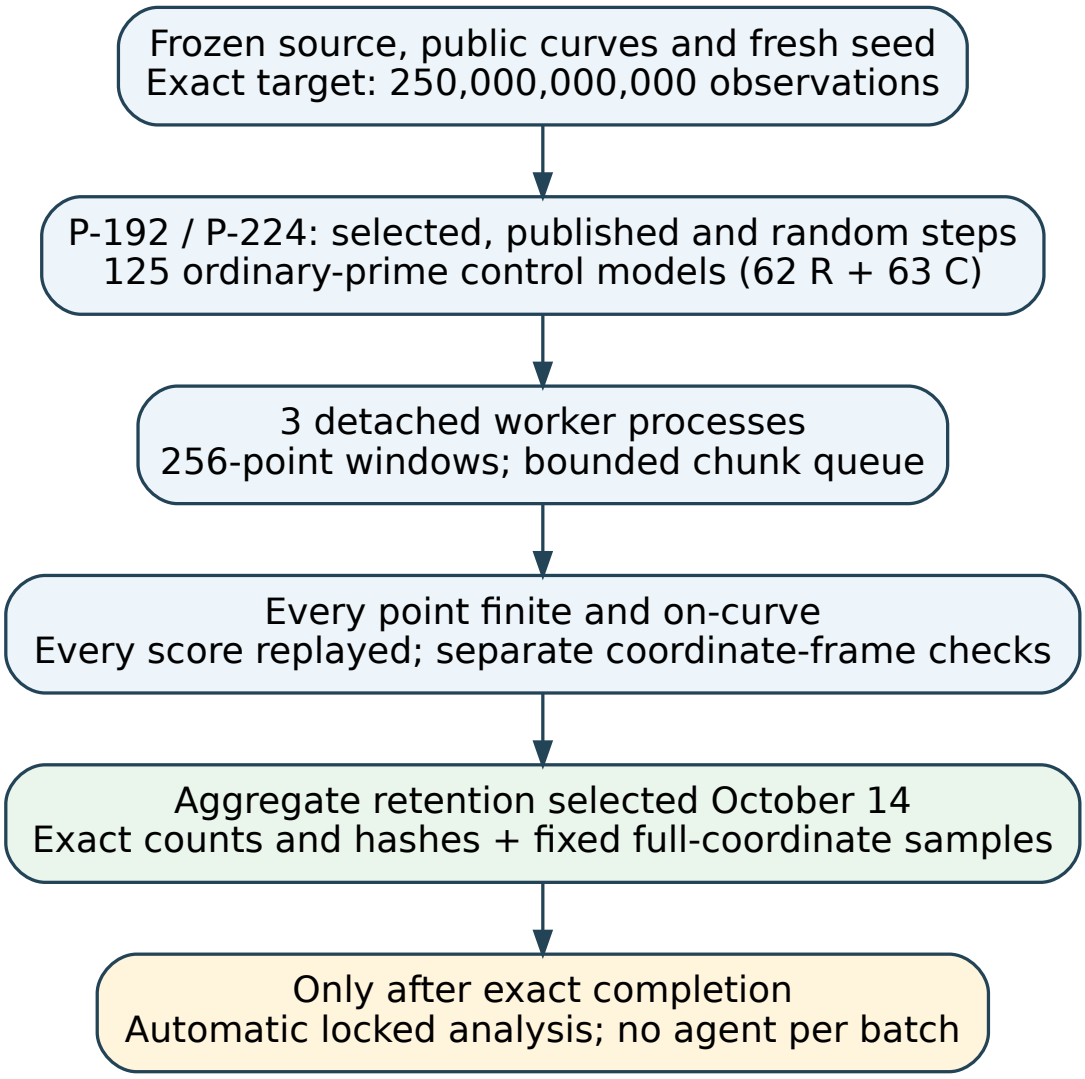

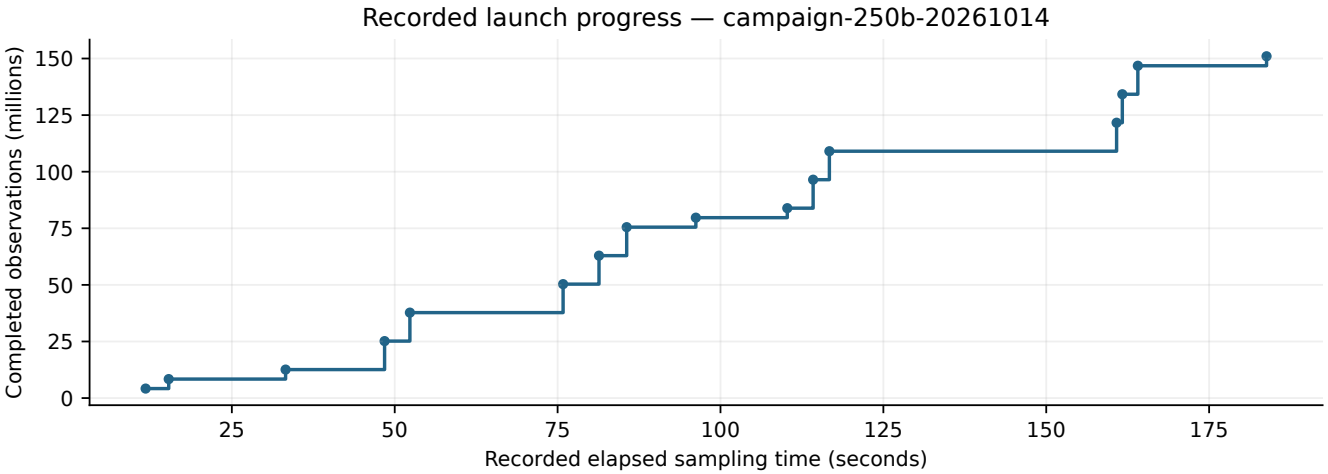

Graph: exact recorded counters, not population-rate estimates; no uncertainty interval applies. Source: publication/launch-20261014/progress-series.json. Within-window points are dependent; named cohorts share starts.

Fixed scope, retention and resumption

Population	Planned windows/arm or model	Arms/models
P-192 motif	44,389,200	3
P-224 parity	44,389,200	3
P-224 long run	177,556,850	3
Ordinary-prime controls	1,420,454	125

Derived prospective total: 976,562,500 windows × 256 = 250,000,000,000 point observations; 443,892,000 start draws. Selected / published / fresh random steps; 62 R + 63 C ordinary models. All 125 model preparations checked before launch. Primes reused; coefficients and step x matched; group/step orders unmatched.

Locked inference: 18 exact paired contrasts + 2 approximate field comparisons; alpha 0.0001/20, nominal 99.99% primary family with a field-approximation caveat. Separate 30-rate CP family, nominal at least 99.99% under independent windows; union nominal 99.98%, subject to the approximation. No final inference yet. Analysis runs automatically only after exact completion; no outcome-based stopping.

Aggregate retention: exact counts, hashes and fixed full-coordinate samples. Most individual bit strings/scalars are not retained for later arbitrary ML. Full compact data require approximately 60–70 GB; initial free capacity was about 550 MB, with no configured managed AWS identity. Existing data remain intact.

Measured runtime: 506.049 completed-chunk CPU seconds, excluding in-flight chunks. Three worker processes were active; detached PID 1168, parent PID 1. Python 3.12.14, NumPy 2.3.5, SciPy 1.17.0. Unknown model/usage metadata remain null.

Source pushed before sampling: `068797c4d5d59c9e73531f751e820c9250fe7c9a`. Runtime dirty paths were only the new publication/preparation report; scientific source was unchanged. See runtime.json, compile.json and manifest.json in tasks/solinas_250b/publication/launch-20261014/.

Standalone/resumable command: tasks/solinas_250b/README.md. The identical command and output path resume checksum-sealed chunks. Changed source/config/storage or corruption is refused. The machine must remain active; the run is expected to take days. No agent is needed per batch. This is not an ECDLP-cost or speedup claim.

Evidence: the launch snapshot and sealed chunk records in publication/launch-20261014/; exact public parameters and full curve UIDs in tasks/solinas_250b/inputs.json; protocol/power records in the same task; predecessor PR 2108. Internal sources, read in this session. No canonical mechanism graph or official status changes. Next: fixed completion, locked final analysis and a fresh final-results report.