Larger solver validation: stopped at build

The admitted run reached no confirmation measurements

This is an observations-only report of an operational attempt. Three candidate builds failed the admitted warning policy. Three baseline builds and their fixed self-tests passed. The larger-cohort solver comparison did not start.

Requested: 48 fresh curves, subgroup sizes 12/14/16/18 bits, and 768 public targets. Recorded solver observations: 0. Supervisor status: infrastructure_error; exit 2. The output checker and public-witness verification were not reached.

Module	Build exit	Build(s)	Self-test(s)	Measurement
candidate_0	1	6.591261	not run	not dispatched
candidate_1	1	5.741320	not run	not dispatched
candidate_2	1	7.049104	not run	not dispatched
solvers_0	0	9.412291	0.032652	not dispatched
solvers_1	0	9.307779	0.016191	not dispatched
solvers_2	0	9.076080	0.016734	not dispatched

Exact compiler stop

The admitted command includes -Wshadow and -Werror. The compiler reports that the local b in decoder.hpp:311 shadows the local b at line 300, inside dec::self_test(). All three candidate builds emit the same stderr digest and exit 1.

The driver dispatches its pool only if all six builds and self-tests succeed. Because only three modules were eligible, no measurement command was invoked. The aggregate records six incomplete modules; this is not six compilation failures.

The baseline self-tests each report 269 development checks. These repeat the same fixed development suite; they are neither independent replications nor results from the 768-target confirmation cohort.

Retained evidence and next action

No solver correctness or performance conclusion follows from this attempt

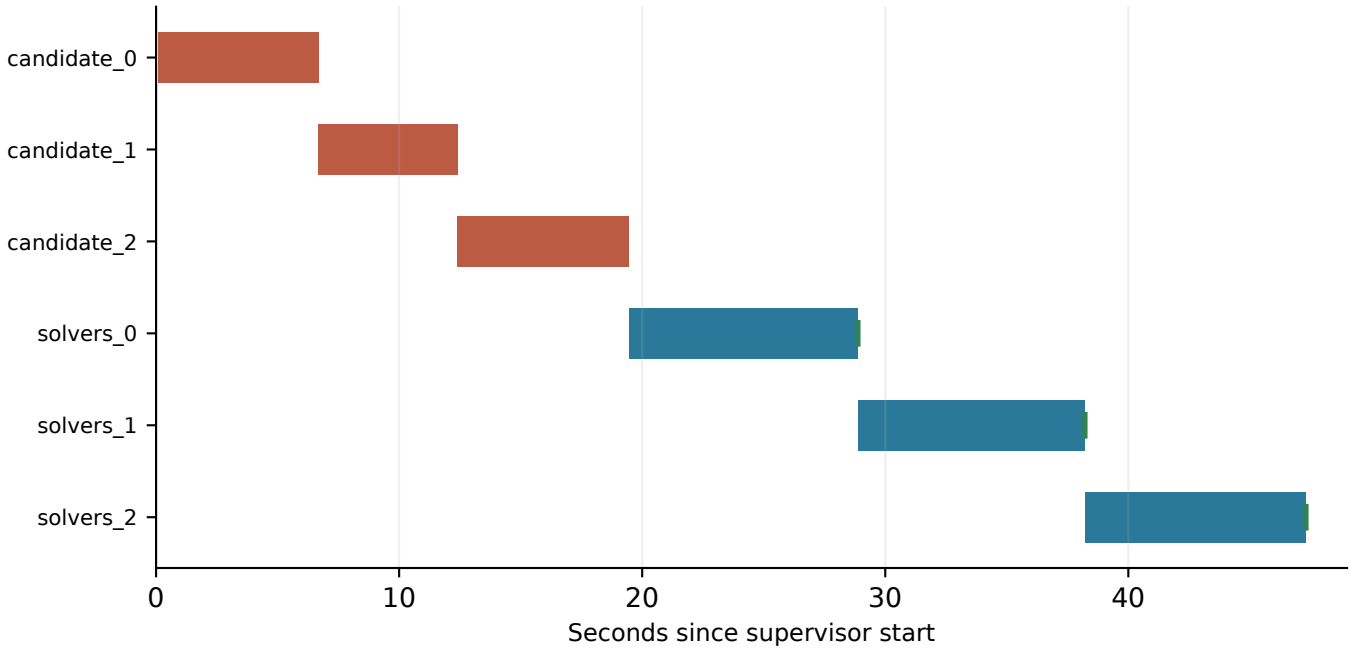

Red: build exit 1; blue: build exit 0; green ticks: self-test exit 0.

The supervisor ran 19:06:32.953559-19:07:20.789667 UTC on 25 October 2026, taking 47.836108 seconds. These are preparation durations, not solver-comparison measurements.

All 2,612 receipt-listed hashes match. There are 2,613 actual files including the receipt, versus 2,883 required paths. The 270 unstarted checker/verification artifacts remain absent. All 2,560 producer slots are empty: 0 rows.

Raw publication: c395a3c6ae40c6647a16a4c1ec4b4541b0bb9a92. Source/claim: af72d170c751d4210eec09d6670af43cd4b2743d. Exact commands, stdout/stderr, digests and failure records remain sealed.

Next: repair the self-test identifier in a new version, test the exact admitted flags, and obtain fresh admission. No confirmation solver outcomes exist to tune on. A later execution gets its own records and report.

Correctness, candidate recovery, fallback dependence and baseline costs remain unanswered here. This same-session report reads no private labels and supplies no independent scientific review.