Run report: EXP-PFDR-bb0d26

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-PFDR-bb0d26/runbook.yaml (sha256 a2d3e9e44fd47d51688a38c415be64744888714bf0d7d72c8522bcc75da5980f)
- Session: experiments/EXP-PFDR-bb0d26/runbook-sessions/RBS-PFDR-bb0d26
- Journal tip: 47f92d815113fa07878922f3ed9d45ddd7a6adbfbcf0920994e22af9efaa009f (18 entries)
- Commit: 98ab449608ff6eb12d053e68a5dc2143fc22d286; tracked files dirty: False
- Started 2026-10-22T12:17:43.195279Z, finished 2026-10-23T22:11:09.400603Z; status completed

Purpose

Measures, per toy cell (bits in 20, 24, 28; rank r in 1, 2, 3; n in 256, 512, 1024, 2048, 4096; 45 cells), the exact count of nontrivial pair-sum coincidences of the x-label set of one rank-r log-GAP base against 20 matched uniformly random label sets and one planted label arithmetic progression on generated prime-order curves, with the kernel self-test and design table, the main census, the analysis of the main block, the replication census of the cells listed in flags.json on fresh curves and the final analysis, each trial validated by check.py. Observations only; check.py re-derives the design table, census completeness, replayed jobs, per-cell statistics and the toy reading from the recorded jobs. Toy curves only. No speedup, no cost against rho and no statement beyond the toy scope is claimed.

Steps

step	status	what it does
S01-trials	passed	Runs run.py then check.py for every trial of trial-plan.json under the supervisor claim and keeps every output file; census-toy-rep runs in every case and writes an empty toy-rep-jobs.jsonl when flags.json lists no cell.

Units

5 valid of 5; invalid: none

trial	run	receipt status	valid
selftest (control)	RUN-PFDR-e44a19	output_validated	True
census-toy-main	RUN-PFDR-4d8c71	output_validated	True
analyze-main	RUN-PFDR-243425	output_validated	True
census-toy-rep	RUN-PFDR-cd03e2	output_validated	True
analyze-final	RUN-PFDR-089edc	output_validated	True

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
calibration_mean_final	derived	z-score	1	count=1, max=-0.26757, min=-0.26757
calibration_n_final	derived	cells	1	count=1, max=44, min=44
calibration_sd_final	derived	z-score	1	count=1, max=0.97357, min=0.97357
census_main_failed_infrastructure	measured	jobs	1	count=1, max=0, sum=0
census_main_jobs	measured	jobs	1	count=1, max=6657, sum=6657
census_rep_failed_infrastructure	measured	jobs	1	count=1, max=0, sum=0
census_rep_jobs	measured	jobs	1	count=1, max=0, sum=0
fraction_valid_final	derived	fraction	1	count=1, max=0.977778, min=0.977778
selftest_cells	measured	cells	1	count=1, max=45, min=45

EXP-PFDR-bb0d26: steps as recorded (fewer than two comparable values to plot)

S01-trials: passed

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: complete
- Control validity: passed
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)