Run report: EXP-PFDR-ca5e01

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-PFDR-ca5e01/runbook.yaml (sha256 003f8e4d525ece0030a06ee53c59726dbd8a9ca8676efd48f25415be43552b34)
- Session: experiments/EXP-PFDR-ca5e01/runbook-sessions/RBS-PFDR-ca5e01
- Journal tip: b82cdf03bef1a0168258fabad0979419f620487c540cd92db3c9906b71c57565 (10 entries)
- Commit: b6dbfa28bbe0e48d328ca4da34269956166355fe; tracked files dirty: False
- Started 2026-10-21T14:23:45.606954Z, finished 2026-10-21T14:52:30.298020Z; status completed

Purpose

Measures, per cell (bits in 16, 18, 20, 22; theta in 0.30..0.65; d in 1, 2, 3; 90 cells), the exact count of nontrivial pair-sum coincidences of one structured symmetric base (window for d = 1, rank-d label GAP for d = 2, 3) against 20 matched uniformly random symmetric null bases and one planted known-log base on prime-order curves, with the kernel self-test, the pilot sizing of curves per cell, the main census, the analysis of the main block, the replication census of flagged cells on fresh curves and the final analysis, each trial validated by check.py. Observations only; check.py re-derives sizing, completeness, replayed jobs, per-cell statistics and the outcome from the recorded jobs. No speedup, no cost against rho and no statement beyond the toy scope is claimed.

Steps

step	status	what it does
S01-trials	passed	Runs run.py then check.py for every trial of trial-plan.json under the supervisor claim and keeps every output file.

Units

6 valid of 6; invalid: none

trial	run	receipt status	valid
selftest (control)	RUN-PFDR-936605	output_validated	True
pilot	RUN-PFDR-93c865	output_validated	True
census-main	RUN-PFDR-eb4733	output_validated	True
analyze-main	RUN-PFDR-c08b70	output_validated	True
census-rep	RUN-PFDR-e15482	output_validated	True
analyze-final	RUN-PFDR-b7c582	output_validated	True

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
calibration_mean_final	derived	z-score	1	count=1, max=-0.116821, min=-0.116821
calibration_n_final	derived	cells	1	count=1, max=79, min=79
calibration_sd_final	derived	z-score	1	count=1, max=0.945989, min=0.945989
census_main_failed_infrastructure	measured	jobs	1	count=1, max=0, sum=0
census_main_jobs	measured	jobs	1	count=1, max=19388, sum=19388
census_rep_failed_infrastructure	measured	jobs	1	count=1, max=0, sum=0
census_rep_jobs	measured	jobs	1	count=1, max=0, sum=0
fraction_valid_final	derived	fraction	1	count=1, max=0.877778, min=0.877778
pilot_cells	measured	cells	1	count=1, max=90, min=90
pilot_cells_ok	measured	cells	1	count=1, max=90, min=90

pilot_curves_planned	derived	curves	1	count=1, max=19388, min=19388

EXP-PFDR-ca5e01: steps as recorded (fewer than two comparable values to plot)

S01-trials: passed

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: complete
- Control validity: passed
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)