Paired targetwise confirmation candidate

2026-10-22. Implementation controls only; no scientific verdict about graph-wide relation yield or ECDLP cost.

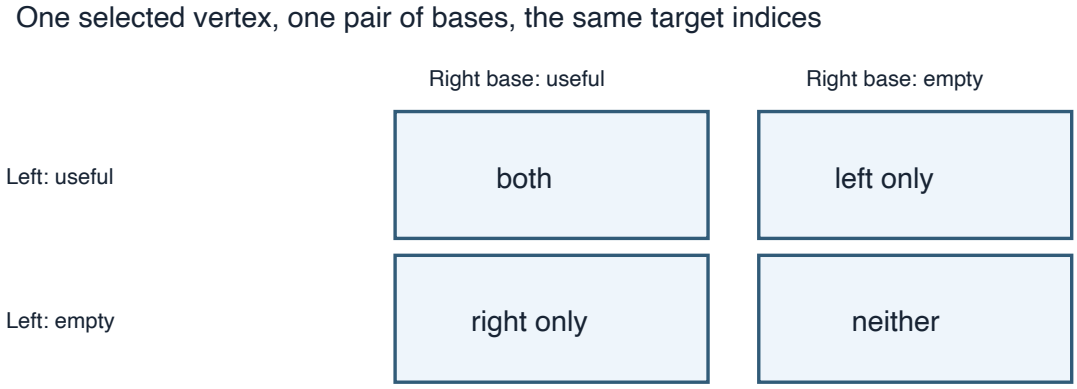

Paired yield difference = (left only – right only) / 32,768; diagram, no graph-wide data.

For every discovery-selected P29/P37 vertex, the candidate binds three fresh confirmation collector streams to supervisor and graph-attempt receipts. It compares the same target indices on low-x, mid-x, and random-x bases. The write-once output retains both/left-only/right-only/neither cells and the paired useful-target difference for each base pair.

Three tests passed. A bounded P37 eight-target fixture yielded 2 useful targets, matching its existing independent rank-verifier count. Synthetic controls rejected changed collector digests, coefficient streams, and target order. These tests do not constitute a full graph comparison.

Next: bind both analysis sources to a Coordinator-approved confirmation plan, collect and validate every fresh arm, then produce paired summaries. Complete same-point ECDLP/rho costs remain a later measured stage.