Prym rank census: implemented, awaiting artifact-store access

EXP-ECDLP-d6d8e8 / H-ECDLP-51f052 / issue #2414. Prospective implementation report, 2026-10-25. No scientific trial or verdict. Six planted software tests pass; independent public-input checks pass.

Question: does full finite branch enumeration yield more useful rank than the 64-branch sampler? Both arms use the same native base and sources 1..64. Publicly exposed columns are removed in both. U=0 is calibration; no stopping when rank becomes full.

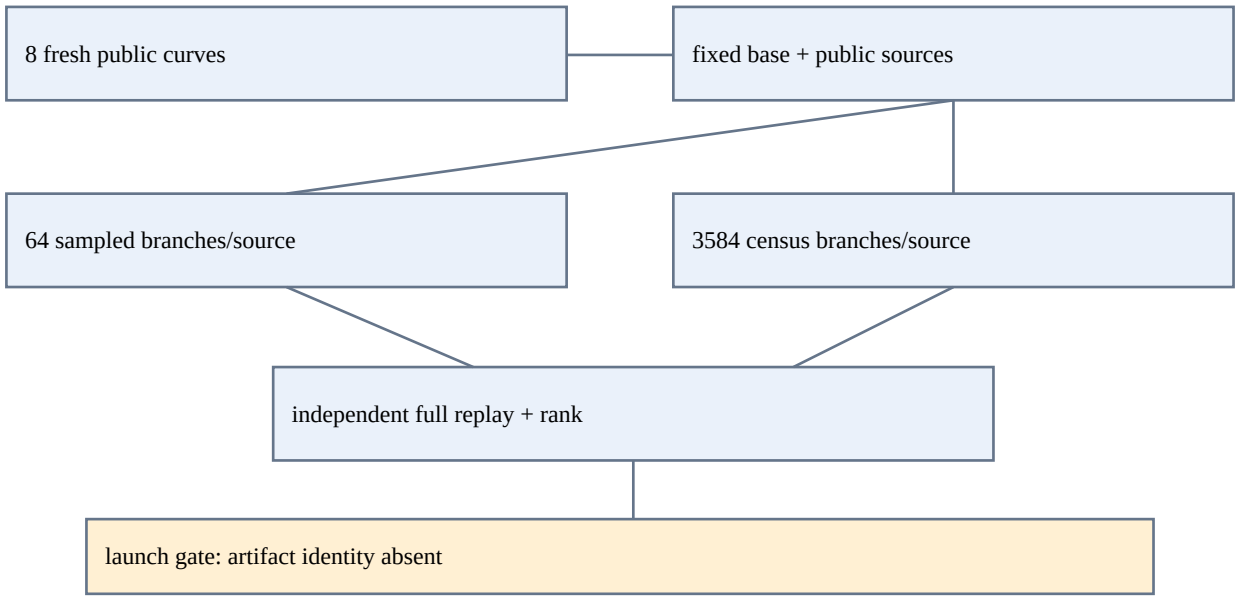

Fixed workload: eight fresh nonanomalous curves at p=101,211,431,4099, A=2, two first-eligible B each. 135 candidate inputs inspected. Derived workload: 8 x 64 x (64+3584) = 1,867,776 attempts. Measured scientific attempts: zero.

Measurements: all statuses, witnesses, rank progression, first full-rank location, exact finite right kernels, partial operation vectors and wall/CPU/RSS diagnostics. Independent checker replays every branch and verifies sample-rowspace containment. Replaying is not a prescribed-computation proof.

Execution gate: repository policy requires large raw data in the canonical S3 artifact store. Normal AWS access returned exit 255: Unable to locate credentials. The runtime has no configured identity. The contract remains review_required. No approval or scientific launch occurred.

Limits: raw counters omit pow internals, Python bookkeeping, byte movement, process-group peak and independent count verification. No complete S, IC gain, asymptotic or deployed-curve claim. New rank results require a native admitted run after access and archival gates; blind recovery additionally needs worker isolation.

Internal sources: README.md, config.json, fixtures.json, input-validation.json, software-tests.txt, infrastructure-status.json; prior RUN-ECDLP-d25f6a. Diagram source: diagram-source.json. Current main merged at ab0609ebe; earlier evidence remains unchanged.